



CREDIT RISK ANALYSIS.

Smitesh Jadhav

PROBLEM STATEMENT

- ▶ Based on the credit history of the customers given, check whether they are capable enough to pay the loan or not. Predict whether the customer will default or not . To manage credit risk by using the past data and deciding whom to give the loan to in the future.

CREDIT RISK ANALYSIS

- ▶ Credit risk or credit default risk associated with a financial transaction is simply the expected loss of that transaction.
- ▶ It can be defined as follows:

$$\text{Credit Risk} = \text{Default Probability} \times \text{Exposure} \times \text{Loss Rate}$$

- ▶ WHERE:

Default Probability is the probability of a debtor reneging on his debt payments.

Exposure is the total amount the lender is supposed to get paid. In most cases, it is simply the amount borrowed by the debtor plus interest payments.

Loss Rate = $1 - \text{Recovery Rate}$, where Recovery Rate is the proportion of the total amount that can be recovered if the debtor defaults. Credit risk analysts analyze each of the determinants of credit risk and try to minimize the aggregate risk faced by an organization.

CREDIT RISK ANALYSIS

- ▶ Credit risk analysis can be thought of as an extension of the credit allocation process.
- ▶ After an individual or business applies to a bank or financial institution for a loan, the lending institution analyzes the potential benefits and costs associated with the loan.
- ▶ Credit risk or credit default risk is a type of risk faced by lenders.
- ▶ Credit risk arises because a debtor can always renege on their debt payments.

DIAGRAMATIC REPRESENTATION



DATASET DESCRIPTION

1. Details about people who applied for loan from June 2007 to December 2015.
2. 73 Variables and 855969 Observations.
3. Dependent Variable: default_ind:- 0 as Not Defaulter
1 as Defaulter

UNDERSTAND THE DATA

jupyter Credit Risk Analysis Last Checkpoint: an hour ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python

```
In [ ]: #IMPORTING LIBRARIES
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
pd.set_option("display.max_rows", None)
pd.set_option("display.max_columns", None)
```

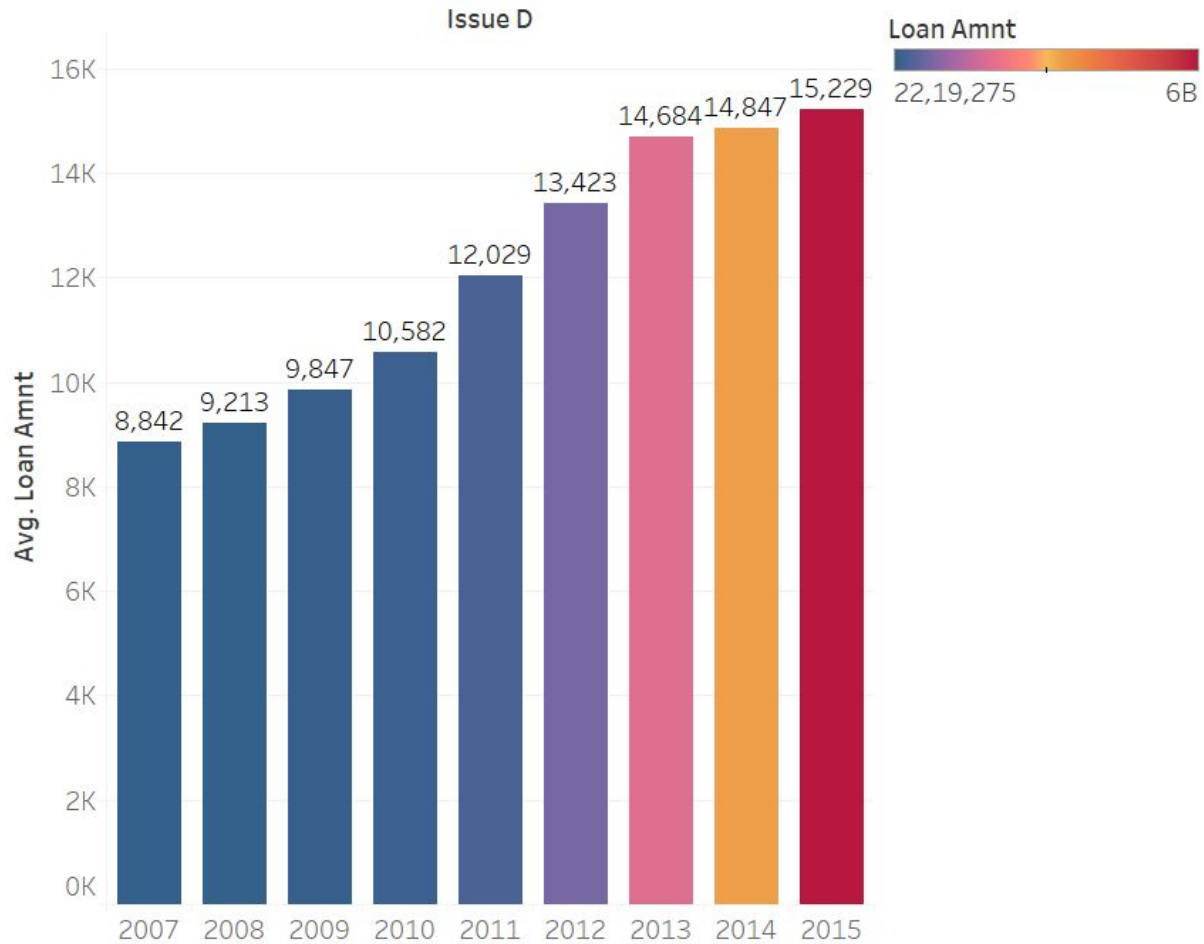
```
In [3]: #LOADING THE DATA & READING IT
Loan_df=pd.read_csv("LoanDefaulter.txt", header=0, index_col=None, sep=" ", delimiter="\t")
```

```
In [4]: Loan_df.head()
```

Out[4]:


	id	member_id	loan_amnt	funded_amnt	funded_amnt_inv	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_ow
0	1077501	1296599	5000.0	5000.0	4975.0	36 months	10.65	162.87	B	B2	NaN	10+ years	
1	1077430	1314167	2500.0	2500.0	2500.0	60 months	15.27	59.83	C	C4	Ryder	< 1 year	
2	1077175	1313524	2400.0	2400.0	2400.0	36 months	15.96	84.33	C	C5	NaN	10+ years	

AVERAGE LOAN AMOUNT GIVEN EVERY YEAR




Average of Loan Amnt for each Issue D Year. Color shows sum of Loan Amnt. The marks are labeled by average of Loan Amnt.

LOAN AMOUNT GIVEN VS RECEIVED EVERY YEAR




Loan Amnt and Total Received Amt for each Issue D Year. Color shows details about Loan Amnt and Total Received Amt.

TOTAL AMOUNT OF LOAN GIVEN EVERY YEAR



The trend of sum of Loan Amnt for Issue D Year. Color shows details about Issue D Year. The marks are labeled by sum of Loan Amnt.

TOTAL AMOUNT RECOVERED EVERY YEAR




The trend of sum of Total Received Amt for Issue D Year. Color shows sum of Total Received Amt. The marks are labeled by sum of Total Received Amt.

Default indicator



Term wise default indicator



DATA PREPROCESSING

- ▶ Removing variables based on domain knowledge.

```
df=df.drop(["id","member_id","desc","zip_code","addr_state","earliest_cr_line"], axis=1)  
df.shape
```

Out[6]: (855969, 67)

```
In [21]: ➤ df=df.drop(["emp_title","title","next_pymnt_d"],axis=1)  
df.shape
```

Out[21]: (855969, 44)

```
In [26]: ➤ df=df.drop(["last_pymnt_d","last_credit_pull_d"],axis=1)  
df.shape
```

Out[26]: (855969, 42)

LIST OF NULLS

```
In [9]: #HANDLING THE MISSING VALUES
df.isnull().sum()
```

```
Out[9]: loan_amnt      0
funded_amnt      0
funded_amnt_inv   0
term             0
int_rate         0
installment      0
grade            0
sub_grade        0
emp_length      43061
home_ownership   0
annual_inc       0
verification_status  0
issue_d          0
pymnt_plan       0
desc            734157
purpose          0
dti              0
delinq_2yrs      0
earliest_cr_line  0
inq_last_6mths   0
mths_since_last_delinq  439812
mths_since_last_record  724785
open_acc         0
pub_rec          0
revol_bal        0
revol_util       446
total_acc        0
initial_list_status  0
out_prncp        0
out_prncp_inv    0
total_pymnt      0
total_pymnt_inv  0
total_rec_prncp  0
total_rec_int    0
total_rec_late_fee  0
recoveries       0
collection_recovery_fee  0
last_pymnt_d     8862
last_pymnt_amnt  0
```

Jupyter LoanDefaulterFinalFile (1) (autosaved)


File Edit View Insert Cell Kernel Widgets Help

Run

```
total_acc      0
initial_list_status  0
out_prncp      0
out_prncp_inv  0
total_pymnt    0
total_pymnt_inv  0
total_rec_prncp  0
total_rec_int  0
total_rec_late_fee  0
recoveries     0
collection_recovery_fee  0
last_pymnt_d   8862
last_pymnt_amnt  0
next_pymnt_d   252971
last_credit_pull_d   50
collections_12_mths_ex_med   56
mths_since_last_major_derog  642830
policy_code     0
application_type  0
annual_inc_joint  855527
dti_joint        855529
verification_status_joint  855527
acc_now_delinq   0
tot_coll_amt    67313
tot_cur_bal     67313
open_acc_6m     842681
open_il_6m      842681
open_il_12m     842681
open_il_24m     842681
mths_since_rcnt_il  843035
total_bal_il    842681
il_util         844360
open_rv_12m     842681
open_rv_24m     842681
max_bal_bc      842681
all_util        842681
total_rev_hi_lim  67313
inq_fi          842681
total_cu_tl     842681
inq_last_12m    842681
default_ind     0
dtype: int64
```


ELIMINATING THE VARIABLES THAT HAVE MORE THAN 50% NA'S

```
In [11]: #PLOTING A GRAPH TO SEE THE MISSING VALUES IN THE DATA
cols = df.isnull().sum()
plt.figure(figsize=(20,10))
cols.plot(kind='bar')
plt.show()
```



In [12]: `#ELIMINATING THE VARIABLES THAT HAVE MORE THAN 50% NA'S`

```
null_cols=[]  
for x in df.columns:  
    if df[x].isnull().sum()>=0.5*len(df):  
        null_cols.append(x)  
print(null_cols)
```

```
['mths_since_last_delinq', 'mths_since_last_record', 'mths_since_last_major_derog', 'annual_inc_joint', 'dti_joint', 'verification_status_joint', 'open_acc_6m', 'open_il_6m', 'open_il_12m', 'open_il_24m', 'mths_since_rcnt_il', 'total_bal_il', 'il_util', 'open_rv_12m', 'open_rv_24m', 'max_bal_bc', 'all_util', 'inq_fi', 'total_cu_tl', 'inq_last_12m']
```

In [13]: `print(len(null_cols))`

```
20
```

In [15]: `df=df.drop(columns=null_cols,axis=1)`
`df.shape`

Out[15]: (855969, 47)

Treating null values

```
: #TREATING THE NULL VALUES OF SOME COLUMNS WITH MEAN BASED ON DOMAIN KNOWLEGDE  
for value in ['revol_util', 'tot_coll_amt', 'tot_cur_bal', 'total_rev_hi_lim']:  
    df[value].fillna(df[value].mean(), inplace=True)
```

```
: #TREATING THE NULL VALUES OF SOME COLUMNS WITH MEDIAN BASED ON DOMAIN KNOWLEGDE  
df['collections_12_mths_ex_med'].fillna(df['collections_12_mths_ex_med'].median(), inplace=True)
```

Checking for any null values

```
In [27]: df.isnull().sum()
```

```
Out[27]: loan_amnt      0
funded_amnt      0
funded_amnt_inv    0
term              0
int_rate          0
installment       0
grade             0
sub_grade         0
emp_length        0
home_ownership    0
annual_inc        0
verification_status 0
issue_d           0
pymnt_plan        0
purpose           0
dti               0
delinq_2yrs       0
inq_last_6mths    0
open_acc          0
pub_rec           0
revol_bal         0
revol_util        0
total_acc         0
initial_list_status 0
out_prncp         0
out_prncp_inv     0
total_pymnt       0
total_pymnt_inv   0
total_rec_prncp   0
total_rec_int     0
total_rec_late_fee 0
recoveries        0
collection_recovery_fee 0
last_pymnt_amnt   0
collections_12_mths_ex_med 0
policy_code       0
application_type  0
acc_now_delinq    0
tot_coll_amt      0
tot_cur_bal       0
total_rev_hi_lim  0
default_ind       0
dtype: int64
```


Converting categorical values into numeric values.

In [36]:

```
#collecting the categorical variables to perform transformation
char_vars=[]
for x in df.columns:
    if df[x].dtype=="object":
        char_vars.append(x)
print(char_vars)
```

```
['term', 'grade', 'home_ownership', 'verification_status', 'pymnt_plan', 'purpose', 'initial_list_status']
```

In [37]:

```
#converting the categorical variables into numeric variables using LabelEncoder()
from sklearn import preprocessing
#create an object
le=preprocessing.LabelEncoder()
for x in char_vars:
    df[x]=le.fit_transform(df[x])
```

Splitting the data.

```
In [38]: #SPLITTING THE DATASET BASED ON THE DATA VARIABLE  
split_date = pd.datetime(2015,5,1)
```

```
In [39]: #SPLITTING INTO TRAIN & TEST  
Train = df.loc[df['issue_d']<=split_date]  
Test = df.loc[df['issue_d']>split_date]
```


```
In [40]: '''ELIMINATING ISSUE_D SINCE WE DON'T NEED IT'''  
Train.drop(columns='issue_d',inplace=True)  
Test.drop(columns='issue_d',inplace=True)
```

```
In [41]: TrainX = Train.drop(columns='default_ind')  
Y_train = Train['default_ind']
```

```
In [42]: TestX = Test.drop(columns='default_ind')  
Y_test = Test['default_ind']
```

Scaling the data

```
In [43]: #SCALING THE DATA
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(TrainX)
X_train= scaler.transform(TrainX)
X_test= scaler.transform(TestX)
```



```
In [40]: df.shape
```

```
Out[40]: (855969, 37)
```



MODEL SELECTION
,TRAINING,TUNNING AND
EVALUATION.

LOGISTIC REGRESSION

LOGISTIC REGRESSION

#MODEL BUILDING USING LOGISTIC REGRESSION

```
from sklearn.linear_model import LogisticRegression
#create the model object
lg=LogisticRegression()
#train the object
lg.fit(X_train,Y_train)
#predict
Y_pred=lg.predict(X_test)
```

#EVALUATION METRICS

```
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
print(confusion_matrix(Y_test,Y_pred))
print(classification_report(Y_test,Y_pred))
print(accuracy_score(Y_test,Y_pred))
```


```
[[256630   50]
 [    63  248]]
```

		precision	recall	f1-score	support
	0	1.00	1.00	1.00	256680
	1	0.83	0.80	0.81	311
accuracy				1.00	256991
macro avg		0.92	0.90	0.91	256991
weighted avg		1.00	1.00	1.00	256991

0.9995602958858482

#PLOT THE ROC CURVE

```
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.show()
```



#ROC CURVE

```
from sklearn import metrics
fpr, tpr, z = metrics.roc_curve(Y_test, Y_pred)
auc = metrics.auc(fpr, tpr)
print(auc)
```

0.8986164288287692

TUNNED LOGISTIC REGRESSION.

```
'''HERE WE CAN SEE THAT 0.73 IS THE BEST THRESHOLD'''
y_pred_class=[]
for value in y_pred_prob[:,1]:
    if value > 0.73:
        y_pred_class.append(1)
    else:
        y_pred_class.append(0)
#print(y_pred_class)
```

#EVALUATION METRICS AFTER TUNNING

```
cfm=confusion_matrix(Y_test, y_pred_class)
print(cfm)
print('Classification report: ')
print(classification_report(Y_test,y_pred_class))
acc= accuracy_score(Y_test, y_pred_class)
print("Accuracy of the model: ", acc)
```

```
[[256646    34]
 [    63   248]]
```

Classification report:


	precision	recall	f1-score	support
0	1.00	1.00	1.00	256680
1	0.88	0.80	0.84	311
accuracy			1.00	256991
macro avg	0.94	0.90	0.92	256991
weighted avg	1.00	1.00	1.00	256991

Accuracy of the model: 0.9996225548754626

#PLOTING THE ROC CURVE

```
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')

plt.show()
```



#ROC CURVE

```
from sklearn import metrics
fpr, tpr, z = metrics.roc_curve(Y_test, Y_pred)
auc = metrics.auc(fpr, tpr)
print(auc)
```

0.8986164288287692

SGD CLASSIFIER.

```
from sklearn.linear_model import SGDClassifier
#create a model
classifier=SGDClassifier()
#fitting training data to the model
classifier.fit(X_train,Y_train)
Y_pred=classifier.predict(X_test)
```

```
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
confusion_matrix = confusion_matrix(Y_test,Y_pred)
print(confusion_matrix)
print("CLASSIFICATION MATRIX:")
print(classification_report(Y_test,Y_pred))
accuracy_score = accuracy_score(Y_test,Y_pred)
print("ACCURACY OF THE MODEL:",accuracy_score)
```


```
[[256672      8]
 [    63    248]]
```

CLASSIFICATION MATRIX:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	256680
1	0.97	0.80	0.87	311
accuracy			1.00	256991
macro avg	0.98	0.90	0.94	256991
weighted avg	1.00	1.00	1.00	256991

ACCURACY OF THE MODEL: 0.999723725733586

```
#PLOTING THE ROC CURVE
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.show()
```



```
#ROC CURVE
from sklearn import metrics
fpr, tpr, z = metrics.roc_curve(Y_test, Y_pred)
auc = metrics.auc(fpr, tpr)
print(auc)
```

0.8986982427605131

DECISION TREE

#MODEL BUILDING USING DECISION TREE

```
from sklearn.tree import DecisionTreeClassifier
```

```
model_DT=DecisionTreeClassifier()
```

##Train the model

```
model_DT.fit(X_train,Y_train)
```

```
Y_pred=model_DT.predict(X_test)
```

#EVALUATION METRICS

```
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
```

```
print(confusion_matrix(Y_test,Y_pred))
```

```
print(classification_report(Y_test,Y_pred))
```

```
print(accuracy_score(Y_test,Y_pred))
```

```
[[ 71217 185463]
 [      1    310]]
```


	precision	recall	f1-score	support
0	1.00	0.28	0.43	256680
1	0.00	1.00	0.00	311
accuracy			0.28	256991
macro avg	0.50	0.64	0.22	256991
weighted avg	1.00	0.28	0.43	256991

```
0.278324921884424
```

#PLOTING THE ROC CURVE

```
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')

plt.show()
```



#ROC CURVE

```
from sklearn import metrics
fpr, tpr, z = metrics.roc_curve(Y_test, Y_pred)
auc = metrics.auc(fpr, tpr)
print(auc)
```

```
0.6371194919343565
```


TUNED DECISION TREE.

#MODEL BUILDING USING PRUNED DECISION TREE AFTER VARIOUS NUMBER OF ITERATIONS

```
from sklearn.tree import DecisionTreeClassifier
dt=DecisionTreeClassifier(criterion="gini",random_state=10,max_depth=50,min_samples_leaf=100)
dt.fit(X_train,Y_train)
Y_pred=dt.predict(X_test)
```

#EVALUATION METRICS

```
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
print(confusion_matrix(Y_test,Y_pred))
print(classification_report(Y_test,Y_pred))
print(accuracy_score(Y_test,Y_pred))
```

```
[[256676    4]
 [   63   248]]
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	256680
1	0.98	0.80	0.88	311


accuracy			1.00	256991
macro avg	0.99	0.90	0.94	256991
weighted avg	1.00	1.00	1.00	256991

0.9997392904809896

#PLOTTING THE ROC CURVE

```
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')

plt.show()
```



#ROC CURVE

```
from sklearn import metrics
fpr, tpr, z = metrics.roc_curve(Y_test, Y_pred)
auc = metrics.auc(fpr, tpr)
print(auc)
```

0.8987060345635363

EXTRA TREE.

#MODEL BUILDING USING EXTRA TREE CLASSIFIER

```
from sklearn.ensemble import ExtraTreesClassifier
model=ExtraTreesClassifier(30,random_state=10,bootstrap=True)
#fit the model on the data and predict the values
model=model.fit(X_train,Y_train)
Y_pred=model.predict(X_test)
```

#EVALUATION METRICS

```
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
print(confusion_matrix(Y_test,Y_pred))
print(classification_report(Y_test,Y_pred))
print(accuracy_score(Y_test,Y_pred))
```

```
[[197451  59229]
 [      22    289]]
```


	precision	recall	f1-score	support
0	1.00	0.77	0.87	256680
1	0.00	0.93	0.01	311
accuracy			0.77	256991
macro avg	0.50	0.85	0.44	256991
weighted avg	1.00	0.77	0.87	256991

0.7694432878972416

#PLOTING THE ROC CURVE

```
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')

plt.show()
```



#ROC CURVE

```
from sklearn import metrics
fpr, tpr, z = metrics.roc_curve(Y_test, Y_pred)
auc = metrics.auc(fpr, tpr)
print(auc)
```

0.8492550497648179

RANDOM FOREST

```
# MODEL BUILDING USING RANDOM FOREST
from sklearn.ensemble import RandomForestClassifier
model_RandomForest=RandomForestClassifier(criterion="entropy",n_estimators=50,random_state=10,
                                          max_depth=50,min_samples_leaf=100)

#fit the model on the data and predict the values
model_RandomForest.fit(X_train,Y_train)
Y_pred=model_RandomForest.predict(X_test)
```


```
#EVALUATION METRICS
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
print(confusion_matrix(Y_test,Y_pred))
print(classification_report(Y_test,Y_pred))
print(accuracy_score(Y_test,Y_pred))
```

```
[[256615    65]
 [    63   248]]
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	256680
1	0.79	0.80	0.79	311
accuracy			1.00	256991
macro avg	0.90	0.90	0.90	256991
weighted avg	1.00	1.00	1.00	256991

```
0.9995019280830846
```

```
#PLOTING THE ROC CURVE
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.show()
```



```
#ROC CURVE
from sklearn import metrics
fpr, tpr, z = metrics.roc_curve(Y_test, Y_pred)
auc = metrics.auc(fpr, tpr)
print(auc)
```

```
0.7702300761592373
```


GRADIENT BOOSTING

#MODEL BUILDING USING GRADIENT BOOST

```
from sklearn.ensemble import GradientBoostingClassifier
model_GradientBoosting=GradientBoostingClassifier(loss="exponential")
#fit the model on the data and predict the values
model_GradientBoosting.fit(X_train,Y_train)
Y_pred=model_GradientBoosting.predict(X_test)
```

#EVALUATION METRICS

```
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
print(confusion_matrix(Y_test,Y_pred))
print(classification_report(Y_test,Y_pred))
print(accuracy_score(Y_test,Y_pred))
```


```
[[245994  10686]
 [    57    254]]
```

	precision	recall	f1-score	support
0	1.00	0.96	0.98	256680
1	0.02	0.82	0.05	311
accuracy			0.96	256991
macro avg	0.51	0.89	0.51	256991
weighted avg	1.00	0.96	0.98	256991

0.9581969796607663

#PLOTTING OF THE ROC CURVE

```
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.show()
```



#ROC CURVE

```
from sklearn import metrics
fpr, tpr, z = metrics.roc_curve(Y_test, Y_pred)
auc = metrics.auc(fpr, tpr)
print(auc)
```

0.8875443268408323

ADA BOOSTING

#MODEL BUILDING USING ADA BOOST

```
from sklearn.ensemble import AdaBoostClassifier
model_AdaBoost=AdaBoostClassifier(n_estimators=50,algorithm="SAMME.R")
#fit the model on the data and predict the values
model_AdaBoost.fit(X_train,Y_train)
Y_pred=model_AdaBoost.predict(X_test)
```

#EVALUATION METRICS

```
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
print(confusion_matrix(Y_test,Y_pred))
print(classification_report(Y_test,Y_pred))
print(accuracy_score(Y_test,Y_pred))
```

```
[[245616 11064]
 [    62   249]]
```


		precision	recall	f1-score	support
	0	1.00	0.96	0.98	256680
	1	0.02	0.80	0.04	311
accuracy				0.96	256991
macro avg		0.51	0.88	0.51	256991
weighted avg		1.00	0.96	0.98	256991

```
0.9567066550968711
```

#PLOT THE ROC CURVE

```
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')

plt.show()
```



#ROC CURVE

```
from sklearn import metrics
fpr, tpr, z = metrics.roc_curve(Y_test, Y_pred)
auc = metrics.auc(fpr, tpr)
print(auc)
```

```
0.8841041205359358
```

COMPARING ALL OUTPUTS

Model Name	Accuracy Score	Type 1 Error	Type 2 Error	Recall Class 0	Recall Class 1	Precision Class 0	Precision Class 1
LOGISTIC	99.95	50	63	1.00	0.80	1.00	0.83
TUNED LOGISTIC	99.96	34	63	1.00	0.80	1.00	0.83
DECISION TREE	27.83	185463	1	0.28	1.00	1.00	0.00
PRUNED DECISION TREE	99.97	4	63	1.00	0.80	1.00	0.98
EXTRA TREES	76.94	59229	22	0.77	0.93	1.00	0.00
RANDOM FOREST	99.95	65	63	1.00	0.80	1.00	0.79
GRADIENT BOOSTING	95.81	10686	57	0.96	0.82	1.00	0.02
ADA	95.67	11064	62	0.96	0.80	1.00	0.02
SGD	99.97	8	63	1.00	0.80	1.00	0.97

CONCLUSION

- ▶ Using analysing techniques one can predict or analyse that a customer applying for the loan will repay the loan or not. So using multiple algorithms we can be able to analyse a defaulter. The best model selected out of all is Pruned Decision Tree with accuracy 99.97%.